

POIR 613: Computational Social Science

Pablo Barberá

School of International Relations
University of Southern California
`pablobarbera.com`

Course website:

pablobarbera.com/POIR613/

Today

1. Project

- ▶ Two-page summary was due on Monday
- ▶ Peer feedback due next Monday
- ▶ See my email for additional details


2. Machine learning

3. Solutions to challenge 5

4. Examples of supervised machine learning

Supervised machine learning

Overview of text as data methods



Outline

- ▶ Supervised learning overview
- ▶ Creating a labeled set and evaluating its reliability
- ▶ Classifier performance metrics
- ▶ One classifier for text
 - ▶ Regularized regression


Supervised machine learning

Goal: classify documents into pre existing categories.

e.g. authors of documents, sentiment of tweets, ideological position of parties based on manifestos, tone of movie reviews...

What we need:

- ▶ Hand-coded dataset (labeled), to be split into:
 - ▶ **Training set**: used to train the classifier
 - ▶ **Validation/Test set**: used to validate the classifier
- ▶ Method to extrapolate from hand coding to unlabeled documents (**classifier**):
 - ▶ Naive Bayes, regularized regression, SVM, K-nearest neighbors, BART, ensemble methods...
- ▶ **Performance metric** to choose best classifier and avoid overfitting: confusion matrix, accuracy, precision, recall...



Basic principles of supervised learning

- ▶ **Generalization**: A classifier or a regression algorithm learns to correctly predict output from given inputs not only in previously seen samples but also in previously unseen samples
- ▶ **Overfitting**: A classifier or a regression algorithm learns to correctly predict output from given inputs in previously seen samples but fails to do so in previously unseen samples. This causes poor prediction/generalization.

Supervised v. unsupervised methods compared


- ▶ The **goal** (in text analysis) is to differentiate *documents* from one another, treating them as “bags of words”
- ▶ Different approaches:
 - ▶ *Supervised methods* require a **training set** that exemplify contrasting **classes**, identified by the researcher
 - ▶ *Unsupervised methods* scale documents based on patterns of similarity from the term-document matrix, without requiring a training step
- ▶ Relative **advantage** of supervised methods:
You already know the dimension being scaled, because you set it in the training stage
- ▶ Relative **disadvantage** of supervised methods:
You *must* already know the dimension being scaled, because you have to feed it good sample documents in the training stage

Supervised learning v. dictionary methods

- ▶ Dictionary methods:
 - ▶ Advantage: **not corpus-specific**, cost to apply to a new corpus is trivial
 - ▶ Disadvantage: **not corpus-specific**, so performance on a new corpus is unknown (domain shift)
- ▶ Supervised learning can be conceptualized as a generalization of dictionary methods, where features associated with each categories (and their relative weight) are **learned from the data**
- ▶ By construction, they will **outperform dictionary methods** in classification tasks, as long as training sample is large enough

Dictionaries vs supervised learning


Lexicons' Accuracy in Document Classification
Compared to Machine-Learning Approach



Source: González-Bailón and Paltoglou (2015)

Dictionaries vs supervised learning

Application: sentiment analysis of NYTimes articles



Source: Barberá et al (2019)

Outline

- ▶ Supervised learning overview
- ▶ Creating a labeled set and evaluating its reliability
- ▶ Classifier performance metrics
- ▶ One classifier for text
 - ▶ Regularized regression

Creating a labeled set

How do we obtain a **labeled set**?

- ▶ External sources of annotation

- ▶ Disputed authorship of Federalist papers estimated based on known authors of other documents
- ▶ Party labels for election manifestos
- ▶ Legislative proposals by think tanks (text reuse)

- ▶ Expert annotation

- ▶ “Canonical” dataset in Comparative Manifesto Project
- ▶ In most projects, undergraduate students (expertise comes from training)

- ▶ Crowd-sourced coding

- ▶ **Wisdom of crowds**: aggregated judgments of non-experts converge to judgments of experts at much lower cost (Benoit et al, 2016)
- ▶ Easy to implement with FigureEight or MTurk

Code the Content of a Sample of Tweets

Instructions ▾

In this job, you will be presented with tweets about the recent protests related to race and law enforcement in the U.S.

You will have to read the tweet and answer a set of questions about its content.

Read the tweet below paying close attention to detail:

Tweet ID: 447



El Cid

@JohnGalt2112

 Follow

[#BlackLivesMatter](#) don't matter unless they are taken by a white cop.

4:23 PM - 13 Dec 2014




Is this tweet related to the ongoing debate about law enforcement and race in the United States?

- ☐ Yes
- ☐ No
- ☐ Don't Know


Crowd-sourced text analysis (Benoit et al, 2016 APSR)

FIGURE 3. Expert and Crowd-sourced Estimates of Economic and Social Policy Positions



Crowd-sourced text analysis (Benoit et al, 2016 APSR)

FIGURE 5. Standard Errors of Manifesto-level Policy Estimates as a Function of the Number of Workers, for the Oversampled 1987 and 1997 Manifestos



Note: Each point is the bootstrapped standard deviation of the mean of means aggregate manifesto scores, computed from sentence-level random n subsamples from the codes.

Evaluating the quality of a labeled set

Measures of agreement:

- ▶ **Percent agreement** Very simple:
(number of agreeing ratings) / (total ratings) * 100%
- ▶ **Correlation**
 - ▶ (usually) Pearson's r , aka product-moment correlation
 - ▶ Formula: $r_{AB} = \frac{1}{n-1} \sum_{i=1}^n \left(\frac{A_i - \bar{A}}{s_A} \right) \left(\frac{B_i - \bar{B}}{s_B} \right)$
 - ▶ May also be ordinal, such as Spearman's rho or Kendall's tau-b
 - ▶ Range is [0,1]
- ▶ **Agreement measures**
 - ▶ Take into account not only observed agreement, but also *agreement that would have occurred by chance*
 - ▶ **Cohen's κ** is most common
 - ▶ **Krippendorff's α** is a generalization of Cohen's κ
 - ▶ Both range from [0,1]

Outline

- ▶ Supervised learning overview
- ▶ Creating a labeled set and evaluating its reliability
- ▶ Classifier performance metrics
- ▶ One classifier for text
 - ▶ Regularized regression

Computing performance

Binary outcome variables:


	Actual value	
Classification	Ham	Spam
Ham	True negative	False negative
Spam	False positive	True positive

Confusion matrix:

- ▶ True negatives and true positives are **correct** predictions (to maximize)
- ▶ False positives and false negatives are **incorrect** predictions (to minimize)

Computing performance

Type I error
(false positive)



Type II error
(false negative)



Computing performance: an example

- Performance **metrics**

	Actual value	
	Ham	Spam
Classification		
Ham	True negative	False negative
Spam	False positive	True positive

- **Accuracy**: correct predictions / total of predictions
 - % of units that are correctly predicted
- **Precision** for positive labels: (true positive) / (false positive + true positive)
 - % of units predicted to be positive that are indeed positive
- **Recall** for positive labels: (true positive) / (true positive + false negative)
 - % of units that are positive and are predicted as such

Computing performance: an example

Classification	Actual value	
	Ham	Spam
Ham	600	300
Spam	200	900

$$\text{Accuracy} = (600 + 900) / (600 + 900 + 200 + 300) \\ = 0.75$$

75% of all emails are correctly classified

Computing performance: an example

Classification	Actual value	
	Ham	Spam
Ham	600	300
Spam	200	900

$$\text{Precision} = (900) / (900 + 200) \\ = 0.82$$

83% of all emails predicted to be spam are indeed **spam**



Computing performance: an example

Classification	Actual value	
	Ham	Spam
Ham	600	300
Spam	200	900

$$\text{Recall} = (900) / (900 + 300) \\ = 0.75$$

75% of all **spam** emails are correctly classified

Computing performance: an example

Classification	Actual value		
	Ham	Spam	
Ham	1700	50	Total emails: Ham = 1850 Spam = 150
Spam	150	100	

Accuracy = $(1700+100) / (1700+50+150+100) = 0.90$

Precision (spam) = $(100) / (150+100) = 0.40$

Recall (spam) = $(100) / (50+100) = 0.67$

Accuracy can be misleadingly high!

Imagine extreme scenario: we classify everything as ham –
accuracy would be 92.5%



The trade-off between precision and recall

Two extreme scenarios (but same underlying data):

1) Model predicts *always* spam

	Actual value	
Classif.	Ham	Spam
Ham	0	0
Spam	800	1200

Accuracy = $1200 / 2000 = 0.60$

Precision (spam) = $1200 / (800 + 1200) = 0.60$

Recall (spam) = $1200 / 1200 = 1.00$

High recall but low precision.

2) Model predicts (*almost*) *always* ham
(e.g. only emails with 10+ links as spam)

	Actual value	
Classif.	Ham	Spam
Ham	800	1190
Spam	0	10

Accuracy = $(800 + 10) / 2000 = 0.40$


Precision (spam) = $10 / 10 = 1.0$

Recall (spam) = $10 / 1200 = 0.01$

High precision but low recall.

Measuring performance

- ▶ Classifier is trained to **maximize in-sample performance**
- ▶ But generally we want to apply method to **new data**
- ▶ Danger: **overfitting**



- ▶ Model is too complex, describes noise rather than signal
 - ▶ Focus on features that perform well in labeled data but may not generalize (e.g. “inflation” in 1980s)
 - ▶ In-sample performance better than **out-of-sample performance**
- ▶ Solutions?
 - ▶ Randomly split dataset into training and test set
 - ▶ Cross-validation

Cross-validation

Intuition:

- ▶ Create K training and test sets (“folds”) within training set.
- ▶ For each k in K, run classifier and estimate performance in test set within fold.
- ▶ Choose best classifier based on cross-validated performance



Outline

- ▶ Supervised learning overview
- ▶ Creating a labeled set and evaluating its reliability
- ▶ Classifier performance metrics
- ▶ One classifier for text
 - ▶ Regularized regression

Types of classifiers

General thoughts:

- ▶ Trade-off between accuracy and interpretability
- ▶ Parameters need to be cross-validated

Frequently used classifiers:

- ▶ Naive Bayes
- ▶ Regularized regression
- ▶ SVM
- ▶ Others: k-nearest neighbors, tree-based methods, etc.
- ▶ Ensemble methods

Regularized regression

Assume we have:

- ▶ $i = 1, 2, \dots, N$ documents
- ▶ Each document i is in class $y_i = 0$ or $y_i = 1$
- ▶ $j = 1, 2, \dots, J$ unique features
- ▶ And x_{ij} as the count of feature j in document i

We could build a linear regression model as a classifier, using the values of $\beta_0, \beta_1, \dots, \beta_J$ that minimize:

$$RSS = \sum_{i=1}^N \left(y_i - \beta_0 - \sum_{j=1}^J \beta_j x_{ij} \right)^2$$

But can we?

- ▶ If $J > N$, OLS does not have a unique solution
- ▶ Even with $N > J$, OLS has low bias/high variance (overfitting)

Regularized regression

What can we do? Add a **penalty for model complexity**, such that we now minimize:

$$\sum_{i=1}^N \left(y_i - \beta_0 - \sum_{j=1}^J \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^J \beta_j^2 \rightarrow \text{ridge regression}$$

or

$$\sum_{i=1}^N \left(y_i - \beta_0 - \sum_{j=1}^J \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^J |\beta_j| \rightarrow \text{lasso regression}$$

where λ is the **penalty parameter** (to be estimated)

Regularized regression

Why the penalty (shrinkage)?

- ▶ Reduces the variance
- ▶ Identifies the model if $J > N$
- ▶ Some coefficients become zero (feature selection)

The penalty can take different forms:

- ▶ **Ridge regression**: $\lambda \sum_{j=1}^J \beta_j^2$ with $\lambda > 0$; and when $\lambda = 0$ becomes OLS
- ▶ **Lasso** $\lambda \sum_{j=1}^J |\beta_j|$ where some coefficients become zero.
- ▶ **Elastic Net**: $\lambda_1 \sum_{j=1}^J \beta_j^2 + \lambda_2 \sum_{j=1}^J |\beta_j|$ (best of both worlds?)

How to find best value of λ ? Cross-validation.

Evaluation: regularized regression is easy to interpret, but often outperformed by more complex methods.